

Docker

What it can do and why you would want to use it

(Some) History

- OS-level virtualization on Linux
 - [chroot](#) (1982)
 - just changes apparent root directory
 - a number of limitations (stacking, jail break)
 - [Linux-VServer](#) (2001)
 - uses security contexts provided by kernel
 - jail mechanism
 - [LXC](#) (Linux Containers, 2008)
 - uses kernel cgroups/namespace isolation
 - [Docker](#) (2013)
 - uses libcontainer, libvirt, LXC, systemd-nspawn
 - containers are standard processes

Docker Components

- Software
 - dockerd: daemon, manages containers
 - docker: CLI for dockerd
- Objects
 - image: read-only template for containers
 - container: environment that runs applications
 - service: for scaling across multiple dockerd (aka swarm)
- Registries
 - repository for images (pull/push)
 - public or private
 - main public repo: [Docker Hub](#)
 - other backends, e.g.: [sonatype nexus repository](#)

Source: [Wikipedia](#)

Installation

- CLI

`sudo apt install docker.io`

- **Docker Desktop** (check terms re commercial use!)

<https://docs.docker.com/desktop/install/linux-install/>

Using images

- Pull image

`docker pull <URL>`

`docker pull hello-world`

- URL format

`[registry-url/]namespace/image[:tag]`

- "Official" images have no namespace (implicit "library"), hence only "hello-world"

Manage images

- Sub-command "image"
- list images: `docker image ls`
- inspect image: `docker image inspect ...`
- delete image(s): `docker image rm ...`
- remove unused images: `docker image prune ...`

Using images (2)

- Spin up container from image

`docker run <URL>`

`docker run hello-world`

- Output

Hello from Docker!

This message shows that your installation appears to be working correctly.

...

Using images (3)

- Two types of volumes
 - named volumes: persistent, eg for databases
 - mapped directories
 - v/--volume or --mount (greater control than -v)
 - v HOSTDIR:CONTAINERDIR

Managing Containers

- Sub-command "container"
- List running containers: `docker container ls`
- List all containers: `docker container ls -a`
- Start/Stop container: `docker container start/stop ...`
- Remove container(s): `docker container rm ...`

Create images

- Image specification: Dockerfile
- Basic commands:
 - FROM: the base image
 - RUN: executes command (use)
 - ENV: set persistent environment variable
 - COPY: copy files from host into image (from within context)
 - WORKDIR: sets the current working dir, creates it automatically
- Each command is a layer
- Docker detects changes in Dockerfile by hashing the commands (hence use versions!)
- Build image with sub-command "build" (same dir as Dockerfile, aka context)
docker build -t TAGNAME .

Clean up

- stop all containers
`docker stop $(docker ps -a -q)`
- remove all containers
`docker rm $(docker ps -a -q)`
- purging all unused or dangling images, containers, volumes, and networks:
`docker system prune`
- you can be even more aggressive when adding the `-a` flag:
`docker system prune -a`

Why use it?

- Isolating applications
- Legacy software with outdated libraries
- Conflicting library requirements (eg CUDA)
- Easily deploy complex applications
- Version applications (easily switch versions)

If you want to know more...

- Introduction for Data Scientists

<https://www.data-mining.co.nz/docker-for-data-scientists/>

- Docker reference

<https://docs.docker.com/engine/reference/builder/>